

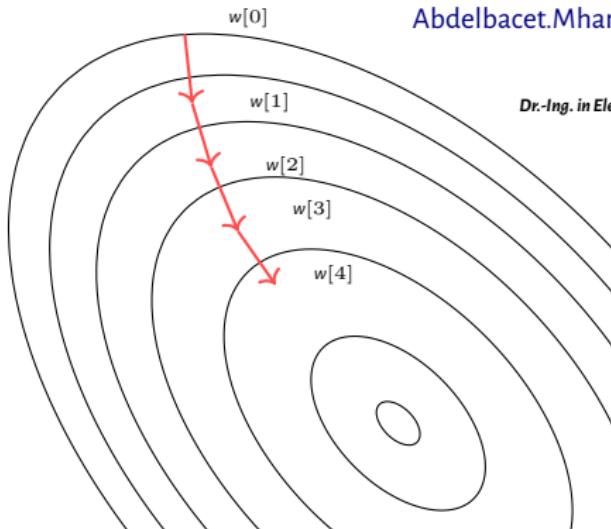
An Introduction to Machine Learning Sorcery

(AN EARLY DRAFT)¹

Abdelbacet Mhamdi

Abdelbacet.Mhamdi@bizerte.r-iset.tn

Dr.-Ing. in Electrical Engineering



“Computers are able to see, hear and learn.
Welcome to the future.”

Dave Waters

“This is nothing. In a few years, that bot will move
so fast you'll need a strobe light to see it.
Sweet dreams...”

Elon Musk

“Machine intelligence is the last invention
that humanity will ever need to make.”

Nick Bostrom

¹Available @ <https://github.com/a-mhamdi/isetbz/>



Disclaimer

- This document features some material gathered from multiple online sources.
- Please note no copyright infringement is intended, and I do not own nor claim to own any of the original material. They are used for educational purposes only.
- I have included links solely as a convenience to the reader. Some links within these slides may lead to other websites, including those operated and maintained by third parties. The presence of such a link does not imply a responsibility for the linked site or an endorsement of the linked site, its operator, or its contents.

- 1 An overview
- 2 Supervised Learning
- 3 Unsupervised Learning
- 4 Deep Learning

Next...



1 An overview

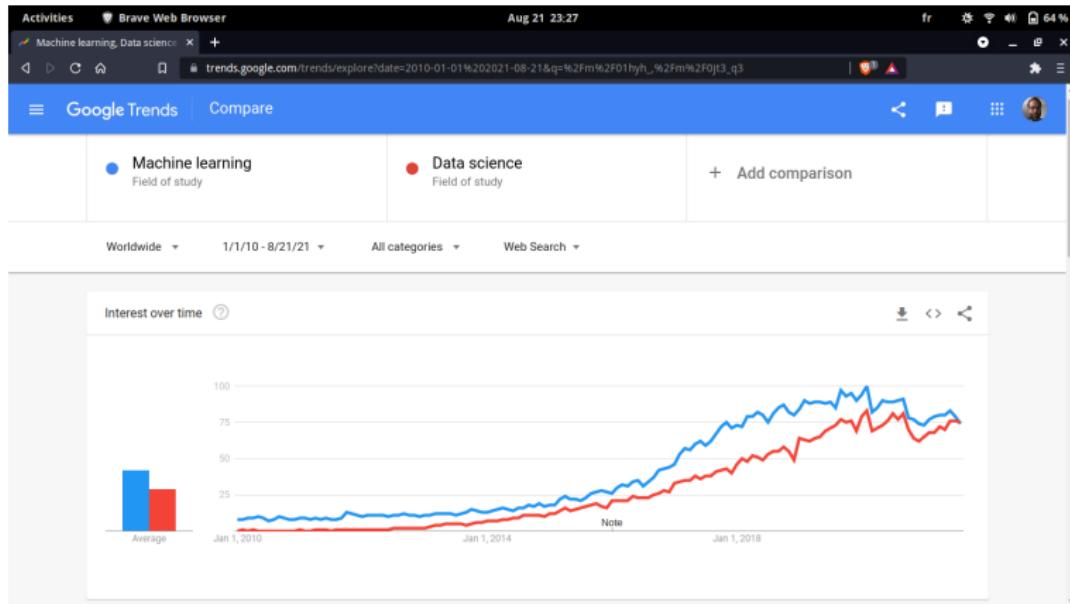
2 Supervised Learning

3 Unsupervised Learning

4 Deep Learning

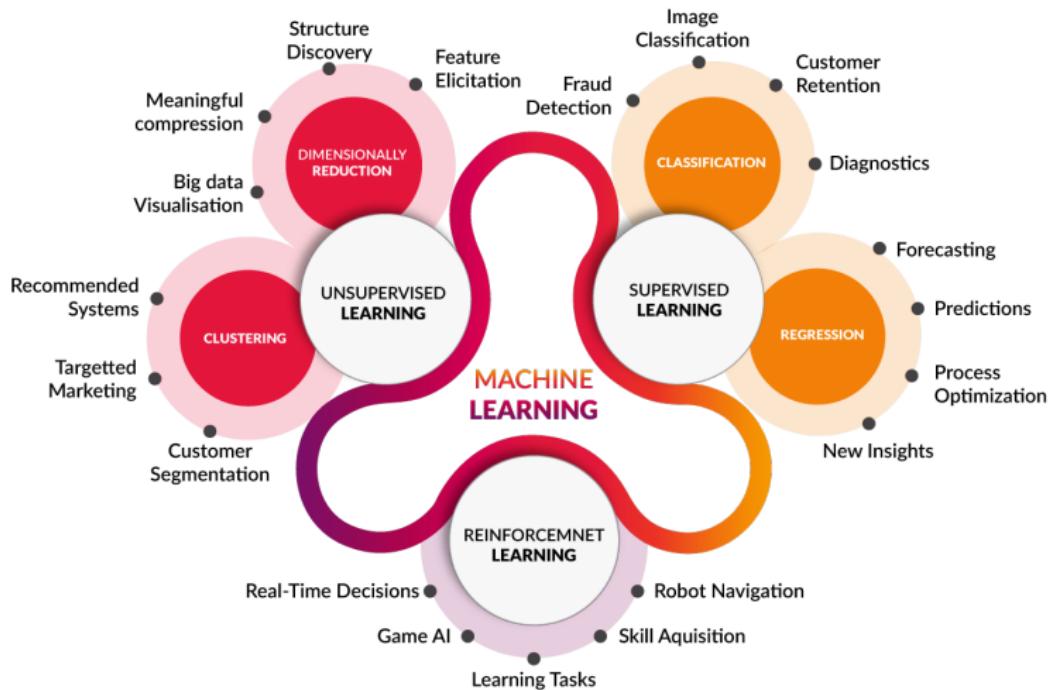
Top uses



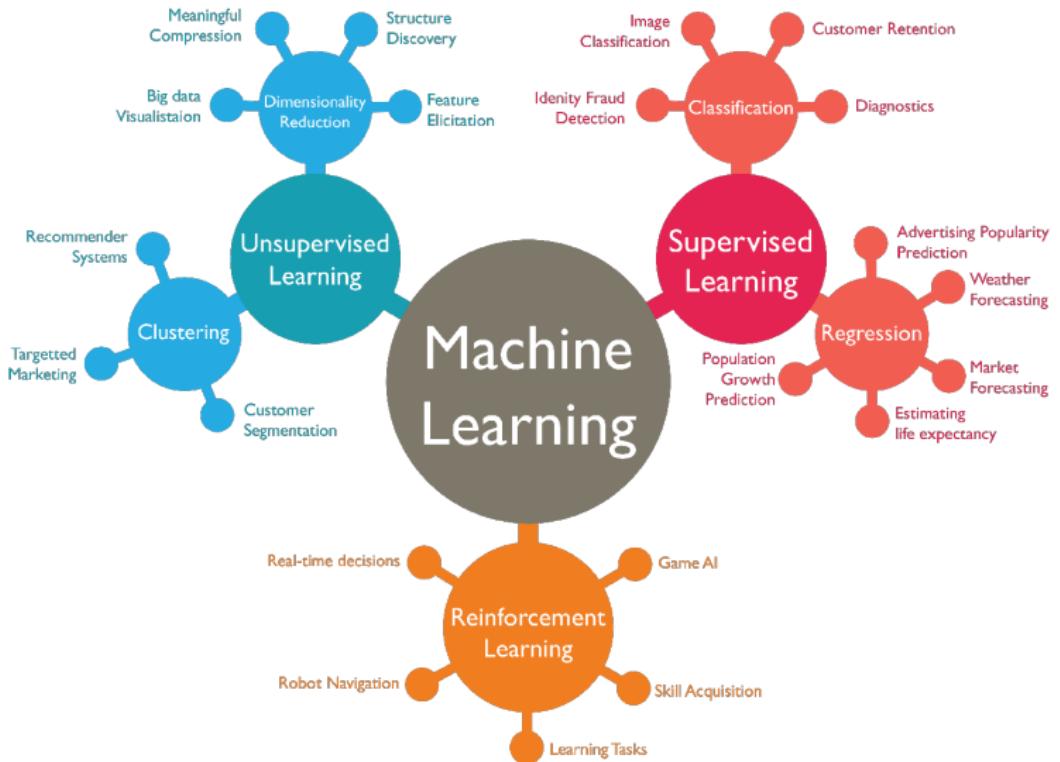


“Numbers represent search interest relative to the highest point on the chart for the given region and time.

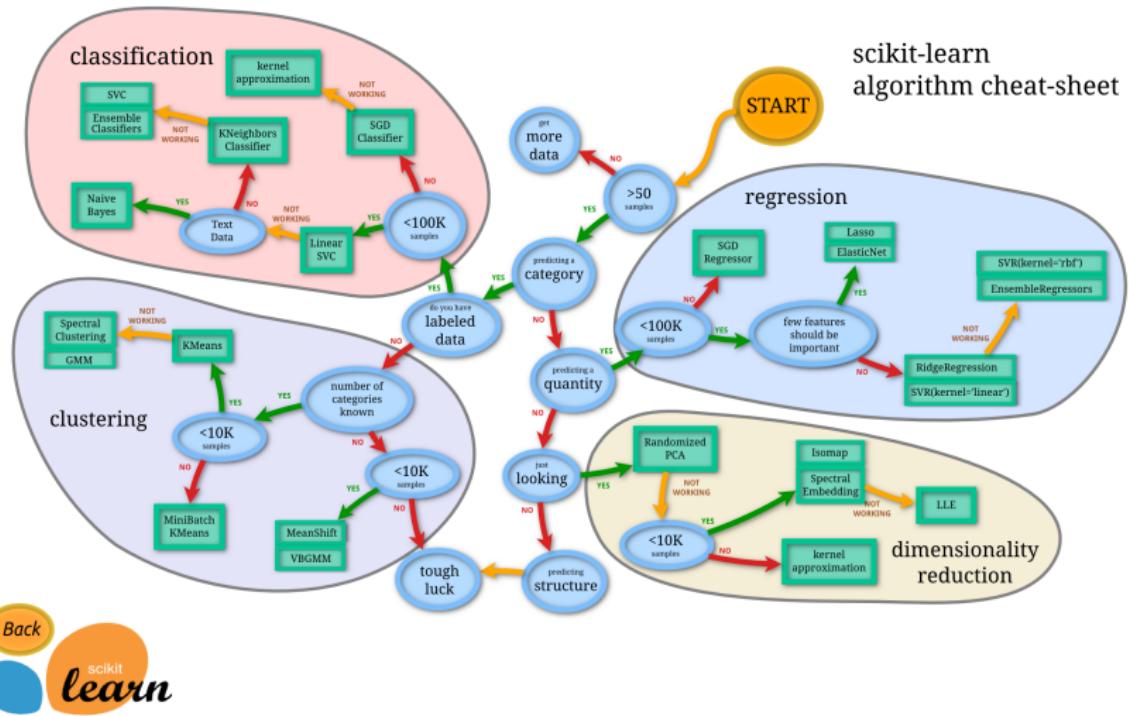
- A value of 100 is the peak popularity for the term;
- A value of 50 means that the term is half as popular;
- A score of 0 means there was not enough data for this term.”



<https://www.cognub.com/index.php/cognitive-platform/>



<https://vitalflux.com/great-mind-maps-for-learning-machine-learning/>



https://scikit-learn.org/stable/tutorial/machine_learning_map/index.html

Regression | Classification | Clustering

<https://github.com/MathWorks-Teaching-Resources/Machine-Learning-for-Regression>



Machine Learning Definition

Arthur Samuel (1959)

Machine Learning: Field of study that gives computers the ability to learn without being explicitly programmed.

Tom Mitchell (1998)

Well-posed Learning Problem: A computer is said to learn from experience \mathcal{E} with respect to some task \mathcal{T} and some performance measure \mathcal{P} , if its performance on \mathcal{T} , as measured by \mathcal{P} , improves with experience \mathcal{E} .

Machine Learning Definition

Arthur Samuel (1959)

Machine Learning: Field of study that gives computers the ability to learn without being explicitly programmed.

Tom Mitchell (1998)

Well-posed Learning Problem: A computer is said to learn from experience \mathcal{E} with respect to some task \mathcal{T} and some performance measure \mathcal{P} , if its performance on \mathcal{T} , as measured by \mathcal{P} , improves with experience \mathcal{E} .

Task #1

Suppose your email program watches which emails you do or do not mark as spam, and based on that learns how to better filter spam. What is the task \mathcal{T} in this setting?

- ① Classifying emails as spam or not spam;
- ② Watching you label emails as spam or not spam;
- ③ The number (or fraction) of emails correctly classified as spam/not spam;
- ④ None of the above-this not a machine learning problem.

Machine Learning Definition

Arthur Samuel (1959)

Machine Learning: Field of study that gives computers the ability to learn without being explicitly programmed.

Tom Mitchell (1998)

Well-posed Learning Problem: A computer is said to learn from experience \mathcal{E} with respect to some task \mathcal{T} and some performance measure \mathcal{P} , if its performance on \mathcal{T} , as measured by \mathcal{P} , improves with experience \mathcal{E} .

Task #1

Suppose your email program watches which emails you do or do not mark as spam, and based on that learns how to better filter spam. What is the task \mathcal{T} in this setting?

- ① Classifying emails as spam or not spam;
- ② Watching you label emails as spam or not spam;
- ③ The number (or fraction) of emails correctly classified as spam/not spam;
- ④ None of the above-this not a machine learning problem.

Next...



1 An overview

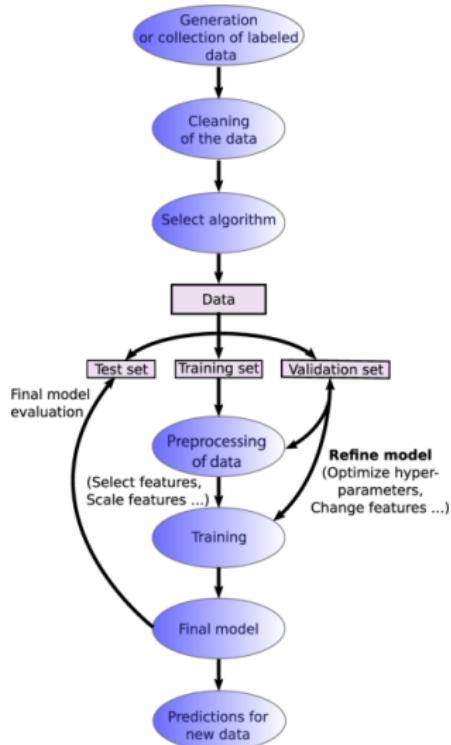
2 Supervised Learning

3 Unsupervised Learning

4 Deep Learning

Overall Methodology

[Sch+19]



Data Preprocessing Template (1/7)

Importing the libraries

```
[1]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt
```

Importing the dataset

```
[2]: df = pd.read_csv('datasets/Data.csv')  
df.head()
```

```
[2]:    Country   Age   Salary Purchased  
0     France  44.0  72000.0      No  
1     Spain   27.0  48000.0     Yes  
2  Germany   30.0  54000.0      No  
3     Spain   38.0  61000.0      No  
4  Germany   40.0       NaN     Yes
```

Extracting independant and dependant variables

```
[3]: X = df.iloc[:, :-1].values  
y = df.iloc[:, -1].values
```

Data Preprocessing Template (2/7)

[4]: `print(X) # Features`

```
[['France' 44.0 72000.0]
 ['Spain' 27.0 48000.0]
 ['Germany' 30.0 54000.0]
 ['Spain' 38.0 61000.0]
 ['Germany' 40.0 nan]
 ['France' 35.0 58000.0]
 ['Spain' nan 52000.0]
 ['France' 48.0 79000.0]
 ['Germany' 50.0 83000.0]
 ['France' 37.0 67000.0]]
```

[5]: `print(y) # Target`

```
['No' 'Yes' 'No' 'No' 'Yes' 'Yes' 'No' 'Yes' 'No' 'Yes']
```

Data Preprocessing Template (3/7)

Imputation transformer for completing missing values

<https://scikit-learn.org/stable/modules/generated/sklearn.impute.SimpleImputer.html>

```
[6]: from sklearn.impute import SimpleImputer
```

```
[7]: imputer = SimpleImputer(missing_values=np.nan, strategy='mean')
imputer.fit(X[:, 1:3])
X[:, 1:3] = imputer.transform(X[:, 1:3])
```

```
[8]: print(X)
```

```
[[ 'France' 44.0 72000.0]
 [ 'Spain' 27.0 48000.0]
 [ 'Germany' 30.0 54000.0]
 [ 'Spain' 38.0 61000.0]
 [ 'Germany' 40.0 63777.77777777778]
 [ 'France' 35.0 58000.0]
 [ 'Spain' 38.77777777777778 52000.0]
 [ 'France' 48.0 79000.0]
 [ 'Germany' 50.0 83000.0]
 [ 'France' 37.0 67000.0]]
```

Data Preprocessing Template (4/7)

How to encode categorical data?

Case of Independent Variable

```
[9]: from sklearn.compose import ColumnTransformer  
      from sklearn.preprocessing import OneHotEncoder
```

```
[10]: ct = ColumnTransformer(transformers=[('encoder', OneHotEncoder(),  
    ↴ [0])], remainder='passthrough')  
X = np.array(ct.fit_transform(X))
```

```
[11]: print(X)
```

```
[[1.0 0.0 0.0 44.0 72000.0]  
 [0.0 0.0 1.0 27.0 48000.0]  
 [0.0 1.0 0.0 30.0 54000.0]  
 [0.0 0.0 1.0 38.0 61000.0]  
 [0.0 1.0 0.0 40.0 63777.7777777778]  
 [1.0 0.0 0.0 35.0 58000.0]  
 [0.0 0.0 1.0 38.77777777777778 52000.0]  
 [1.0 0.0 0.0 48.0 79000.0]  
 [0.0 1.0 0.0 50.0 83000.0]  
 [1.0 0.0 0.0 37.0 67000.0]]
```

Data Preprocessing Template (5/7)

Case of Dependent Variable

```
[12]: from sklearn.preprocessing import LabelEncoder
```

```
[13]: le = LabelEncoder()
y = le.fit_transform(y)
```

```
[14]: print(y)
```

```
[0 1 0 0 1 1 0 1 0 1]
```

Splitting the dataset into Training set and Test set

```
[15]: from sklearn.model_selection import train_test_split
```

https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html

```
[16]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
→2, random_state=1)
```

```
[17]: print(X_train)
```

Data Preprocessing Template (6/7)

```
[[0.0 0.0 1.0 38.77777777777778 52000.0]
 [0.0 1.0 0.0 40.0 63777.7777777778]
 [1.0 0.0 0.0 44.0 72000.0]
 [0.0 0.0 1.0 38.0 61000.0]
 [0.0 0.0 1.0 27.0 48000.0]
 [1.0 0.0 0.0 48.0 79000.0]
 [0.0 1.0 0.0 50.0 83000.0]
 [1.0 0.0 0.0 35.0 58000.0]]
```

```
[18]: print(X_test)
```

```
[[0.0 1.0 0.0 30.0 54000.0]
 [1.0 0.0 0.0 37.0 67000.0]]
```

```
[19]: print(y_train)
```

```
[0 1 0 0 1 1 0 1]
```

```
[20]: print(y_test)
```

```
[0 1]
```

Data Preprocessing Template (7/7)

Scaling of features

```
[21]: from sklearn.preprocessing import StandardScaler
```

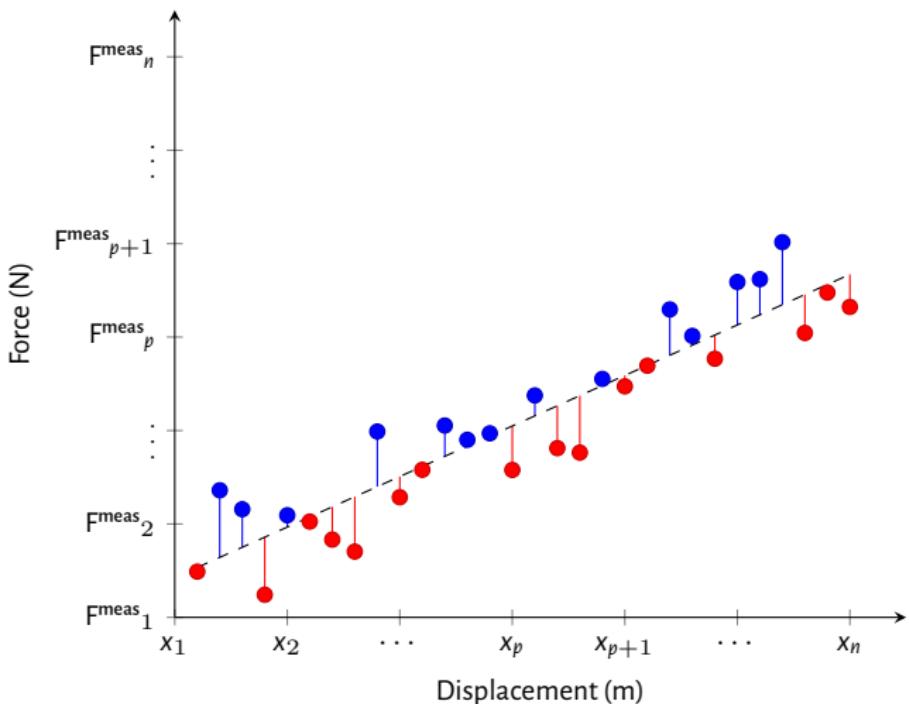
```
[22]: sc = StandardScaler()
X_train[:, 3:] = sc.fit_transform(X_train[:, 3:])
X_test[:, 3:] = sc.transform(X_test[:, 3:])
```

```
[23]: print(X_train)
```

```
[[0.0 0.0 1.0 -0.19159184384578545 -1.0781259408412425]
 [0.0 1.0 0.0 -0.014117293757057777 -0.07013167641635372]
 [1.0 0.0 0.0 0.566708506533324 0.633562432710455]
 [0.0 0.0 1.0 -0.30453019390224867 -0.30786617274297867]
 [0.0 0.0 1.0 -1.9018011447007988 -1.420463615551582]
 [1.0 0.0 0.0 1.1475343068237058 1.232653363453549]
 [0.0 1.0 0.0 1.4379472069688968 1.5749910381638885]
 [1.0 0.0 0.0 -0.7401495441200351 -0.5646194287757332]]
```

```
[24]: print(X_test)
```

```
[[0.0 1.0 0.0 -1.4661817944830124 -0.9069571034860727]
 [1.0 0.0 0.0 -0.44973664397484414 0.2056403393225306]]
```



Consider the example of a spring. Our main goal is to determine the stiffness k of this spring, given some experimental data. The mathematical model (*Hooke's law*):

$$F = kx \quad (1)$$

Restoring force is proportional to displacement.

Table: Measurements of couple (x_i, F^{meas}_i)

x_i	x_1	\dots	x_p	\dots	x_n
F^{meas}_i	F^{meas}_1	\dots	F^{meas}_p	\dots	F^{meas}_n

$$\begin{aligned} F^{\text{meas}}_i &= F_i + \varepsilon_i \\ &= kx_i + \varepsilon_i, \end{aligned} \quad (2)$$

where F_i denotes the unknown real value of the force applied to the spring. In order to estimate the stiffness value k , we can consider the quadratic criterion:

$$\begin{aligned} \mathcal{J} &= \sum_{i=1}^n \varepsilon_i^2 \\ &= \sum_{i=1}^n (F^{\text{meas}}_i - kx_i)^2 \end{aligned}$$

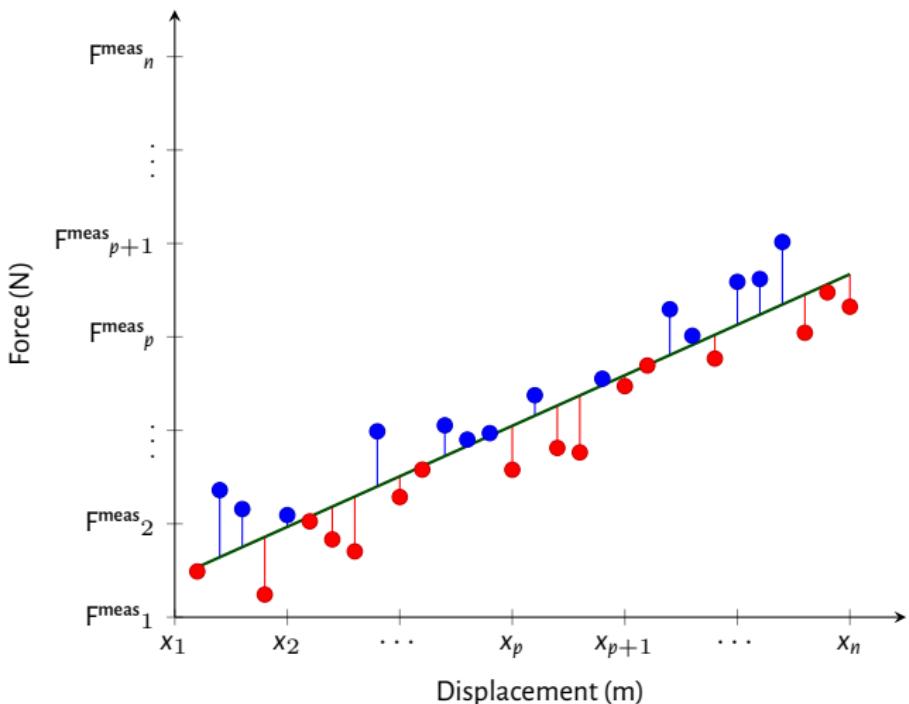
Linear Regression

$$\frac{\partial \mathcal{J}}{\partial k} = 0 \quad (3)$$

$$2 \sum_{i=1}^n (\mathbf{F}^{\text{meas}}_i - kx_i) \sum_{i=1}^n \frac{\partial (\mathbf{F}^{\text{meas}}_i - kx_i)}{\partial k} = 0$$

$$\sum_{i=1}^n (\mathbf{F}^{\text{meas}}_i - kx_i) \sum_{i=1}^n x_i = 0$$

$$\sum_{i=1}^n \mathbf{F}^{\text{meas}}_i x_i = k \sum_{i=1}^n x_i^2 \iff \hat{k} = \frac{\sum_{i=1}^n \mathbf{F}^{\text{meas}}_i x_i}{\sum_{i=1}^n x_i^2}$$



Simple Linear Regression (1/6)

Importing the libraries

```
[1]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt
```

Importing the dataset

```
[2]: df = pd.read_csv('datasets/Salary_Data.csv')  
df.head()
```

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0

```
[3]: X = df.iloc[:, :-1].values  
y = df.iloc[:, -1].values
```

Splitting the dataset into training set and test set

Simple Linear Regression (2/6)

```
[4]: from sklearn.model_selection import train_test_split
```

```
[5]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=1/  
         ↪3, random_state=0)
```

Training the Simple Linear Regression model on the Training set

```
[6]: from sklearn.linear_model import LinearRegression
```

```
[7]: regressor = LinearRegression()  
regressor.fit(X_train, y_train)
```

```
[7]: LinearRegression()
```

Predicting the Test set results

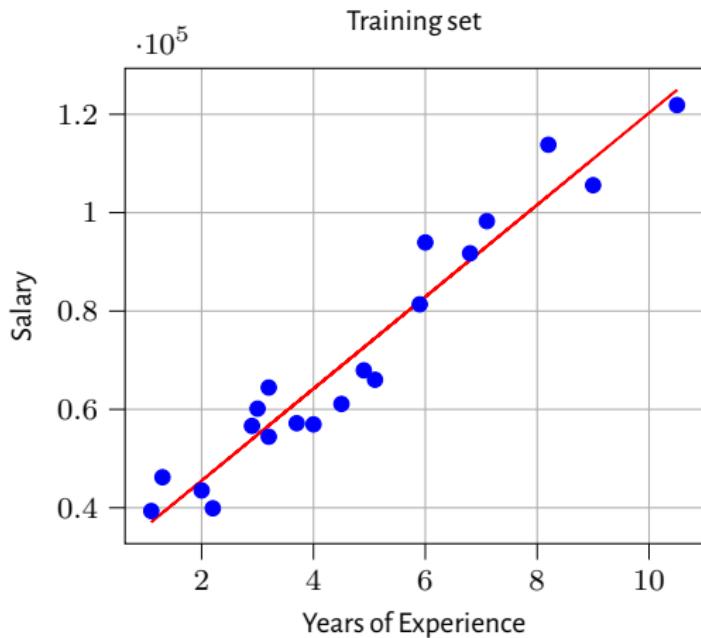
```
[8]: y_pred = regressor.predict(X_test)
```

Visualising the Training set results

Simple Linear Regression (3/6)

```
[9]: plt.scatter(X_train, y_train, color='blue')
plt.plot(X_train, regressor.predict(X_train), color='red')
plt.title('Training set')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.grid()
```

Simple Linear Regression (4/6)

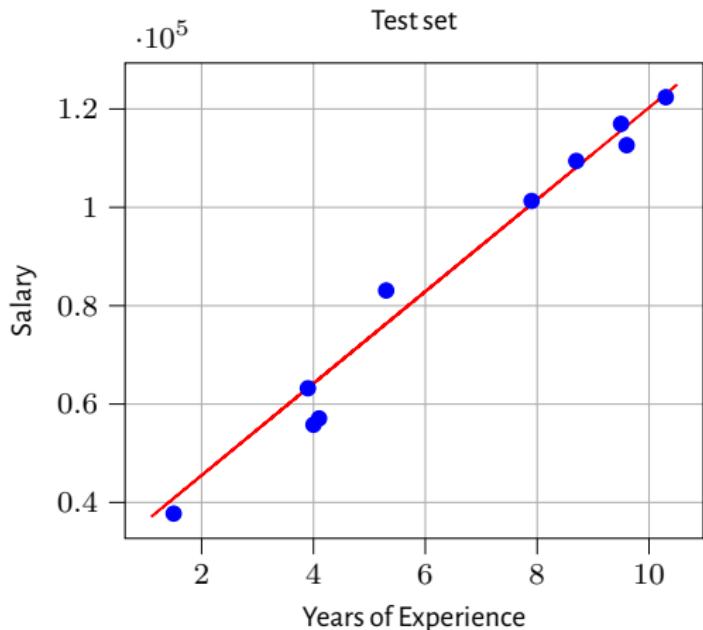


Simple Linear Regression (5/6)

Visualising the Test set results

```
[10]: plt.scatter(X_test, y_test, color='blue')
plt.plot(X_train, regressor.predict(X_train), color='red')
plt.title('Test set')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.grid()
```

Simple Linear Regression (6/6)



This example consists on determining the unknown couple (y_0, v_0) of a mobile solid. We assume that the trajectory is linear. The mathematical model that relates the position y to time t is given by this equation:

$$y = y_0 + v_0 t \quad (4)$$

Table: Measurements of position y

t_i	t_1	\dots	t_p	\dots	t_n
y^{meas}_i	y^{meas}_1	\dots	y^{meas}_p	\dots	y^{meas}_n

$$\begin{aligned} y^{\text{meas}}_i &= y_i + \varepsilon_i \\ &= y_0 + v_0 t_i + \varepsilon_i, \end{aligned} \quad (5)$$

where y_i denotes the unknown real value of the position y at time point t_i .

In order to estimate the values taken by the couple $[y_0, v_0]^T$, we consider the quadratic criterion again, as follows:

$$\begin{aligned} \mathcal{J} &= \sum_{i=1}^n \varepsilon_i^2 \\ &= \varepsilon^T \times \varepsilon \end{aligned}$$

The vector ε is set by $\varepsilon_i, \forall i \geq 1$:

$$\varepsilon = [\varepsilon_1 \quad \cdots \quad \varepsilon_n]^T$$

$$\frac{\partial \mathcal{J}}{\partial \begin{bmatrix} y_0 \\ v_0 \end{bmatrix}} = 0 \quad (6)$$

Multiple Linear Regression (1/4)

Importing the libraries

```
[1]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt
```

Importing the dataset

```
[2]: df = pd.read_csv('datasets/50_Startups.csv')  
df.head()
```

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94

```
[3]: X = df.iloc[:, :-1].values  
y = df.iloc[:, -1].values
```

```
[4]: print(X[:5])
```

Multiple Linear Regression (2/4)

```
[165349.2 136897.8 471784.1 'New York']
[162597.7 151377.59 443898.53 'California']
[153441.51 101145.55 407934.54 'Florida']
[144372.41 118671.85 383199.62 'New York']
[142107.34 91391.77 366168.42 'Florida']]
```

[5]: `print(y[:5])`

```
[192261.83 191792.06 191050.39 182901.99 166187.94]
```

Encoding categorical data

[6]: `from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder`

[7]: `ct = ColumnTransformer(transformers=[('encoder', OneHotEncoder(), [3]), remainder='passthrough')
X = np.array(ct.fit_transform(X))`

[8]: `print(X[:5])`

Multiple Linear Regression (3/4)

```
[[0.0 0.0 1.0 165349.2 136897.8 471784.1]
 [1.0 0.0 0.0 162597.7 151377.59 443898.53]
 [0.0 1.0 0.0 153441.51 101145.55 407934.54]
 [0.0 0.0 1.0 144372.41 118671.85 383199.62]
 [0.0 1.0 0.0 142107.34 91391.77 366168.42]]
```

Splitting the dataset into training set and test set

```
[9]: from sklearn.model_selection import train_test_split
```

```
[10]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
    ↪2, random_state=0)
```

Training the multiple linear regression model on the training set

```
[11]: from sklearn.linear_model import LinearRegression
```

```
[12]: lr = LinearRegression()
lr.fit(X_train, y_train)
```

```
[12]: LinearRegression()
```

Making predictions using the X test set and comparison

Multiple Linear Regression (4/4)

```
[13]: y_pred = lr.predict(X_test)
np.set_printoptions(precision=2)
print(type(y_pred), y_pred.shape)
y_pred = y_pred.reshape(-1,1)
print(type(y_pred), y_pred.shape)
y_test = y_test.reshape(-1,1)
print(np.concatenate((y_pred, y_test), axis=1))
```

```
<class 'numpy.ndarray'> (10,)
<class 'numpy.ndarray'> (10, 1)
[[103015.2 103282.38]
 [132582.28 144259.4 ]
 [132447.74 146121.95]
 [ 71976.1   77798.83]
 [178537.48 191050.39]
 [116161.24 105008.31]
 [ 67851.69  81229.06]
 [ 98791.73  97483.56]
 [113969.44 110352.25]
 [167921.07 166187.94]]
```

Consider the following multivariate equation:

$$y = \theta_1 x_{(1)} + \theta_2 x_{(2)} + \cdots + \theta_m x_{(m)} \quad (7)$$

For a particular single measurement, eq. (7) can be updated as

$$y_k = \theta_1 x_{(1, k)} + \theta_2 x_{(2, k)} + \cdots + \theta_m x_{(m, k)} + \varepsilon_k \quad (8)$$

We denote hereafter by θ the vector $\begin{bmatrix} \theta_1 \\ \theta_2 \\ \vdots \\ \theta_n \end{bmatrix}$. The function y_k becomes:

$$y_k = \underbrace{[x_{(1, k)}, x_{(2, k)}, \dots, x_{(m, k)}]}_{x_k^T} \theta + \varepsilon_k$$

We assume that we have n measurements for y . Then we can transform the previous equation into

$$y = H\theta + \varepsilon,$$

where $y^T = [y_1, y_2, \dots, y_n]$, $X = \begin{bmatrix} x_1^T \\ x_2^T \\ \vdots \\ x_n^T \end{bmatrix}$, and $\varepsilon^T = [\varepsilon_1, \varepsilon_2, \dots, \varepsilon_n]$.

We can consider the mean squared error or quadratic criterion in order to compute the approximated value of θ :

$$\begin{aligned}\mathcal{J} &= \sum_{k=1}^n \varepsilon_k^2 \\ &= \varepsilon^T \varepsilon\end{aligned}$$

The best well estimated value of $\hat{\theta}$ corresponds to the absolute minimum of \mathcal{J} . This leads to calculate the gradient of \mathcal{J} with respect to θ :

$$\frac{\partial \mathcal{J}}{\partial \theta} = \frac{\partial (\varepsilon^T \varepsilon)}{\partial \theta} \quad (9)$$

$$\frac{\partial (\varepsilon^T \varepsilon)}{\partial \theta} = 2 \left(\frac{\partial \varepsilon}{\partial \theta} \right)^T \varepsilon \quad (10)$$

Recall that $\varepsilon = y - X\theta$, the term $\frac{\partial \varepsilon}{\partial \theta}$ hence becomes:

$$\frac{\partial \varepsilon}{\partial \theta} = -X \quad (11)$$

$$\begin{aligned}\frac{\partial J}{\partial \theta} &= 2(-X)^T(y - X\theta) \\ &= 0\end{aligned}$$

The regressor is given by

$$\hat{\theta} = (X^T X)^{-1} X^T y$$



$X^T X$ is not invertible (singular/degenerate)

▼ Redundant Features

Some features are linearly dependant, i.e., \exists some $x_p \propto$ some x_i for instance x_p in feet and x_i in m.

▼ Too many features

Fewer observations compared to the number of features, i.e., $m \geq n$.

- ▲ Delete some features
- ▲ Add extra observations
- ▲ Use regularization

Gradient Descent

$$\theta_i = \theta_i - \underbrace{\alpha}_{\text{LEARNING RATE}} \frac{\partial \mathcal{J}}{\partial \theta_i}$$

Recall that $\mathcal{J} = 1/2n \sum_{k=1}^n (y_k - h_\theta(x_k))^2 \implies \frac{\partial \mathcal{J}}{\partial \theta_i} = 1/n \sum_{k=1}^n (y_k - h_\theta(x_k)) x_{(i, k)}$

$$\theta_i \leftarrow \theta_i - \alpha \frac{1}{n} \sum_{k=1}^n (y_k - h_\theta(x_k)) x_{(i, k)}$$

$$\theta_0 \leftarrow \theta_0 - \alpha \frac{1}{n} \sum_{k=1}^n (y_k - h_\theta(x_k)) x_{(0, k)}$$

$$\theta_1 \leftarrow \theta_1 - \alpha \frac{1}{n} \sum_{k=1}^n (y_k - h_\theta(x_k)) x_{(1, k)}$$

⋮

$$\theta_m \leftarrow \theta_m - \alpha \frac{1}{n} \sum_{k=1}^n (y_k - h_\theta(x_k)) x_{(m, k)}$$

Polynomial Regression (1/8)

Importing the libraries

```
[1]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt
```

Importing the dataset

```
[2]: df = pd.read_csv('datasets/Position_Salaries.csv')  
df.head()
```

```
[2]:
```

	Position	Level	Salary
0	Business Analyst	1	45000
1	Junior Consultant	2	50000
2	Senior Consultant	3	60000
3	Manager	4	80000
4	Country Manager	5	110000

```
[3]: X = df.iloc[:, 1:-1].values  
y = df.iloc[:, -1].values
```

```
[4]: print(type(X), X[:5], sep='\n')
```



Polynomial Regression (2/8)

```
<class 'numpy.ndarray'>
[[1]
 [2]
 [3]
 [4]
 [5]]
```

```
[5]: print(type(y), y[:5], sep='\n')
```

```
<class 'numpy.ndarray'>
[ 45000  50000  60000  80000 110000]
```

Training the linear regression model on the whole dataset

```
[6]: from sklearn.linear_model import LinearRegression
```

Training the linear regression model on the whole dataset

```
[7]: lr_1 = LinearRegression()
lr_1.fit(X, y)
```

```
[7]: LinearRegression()
```

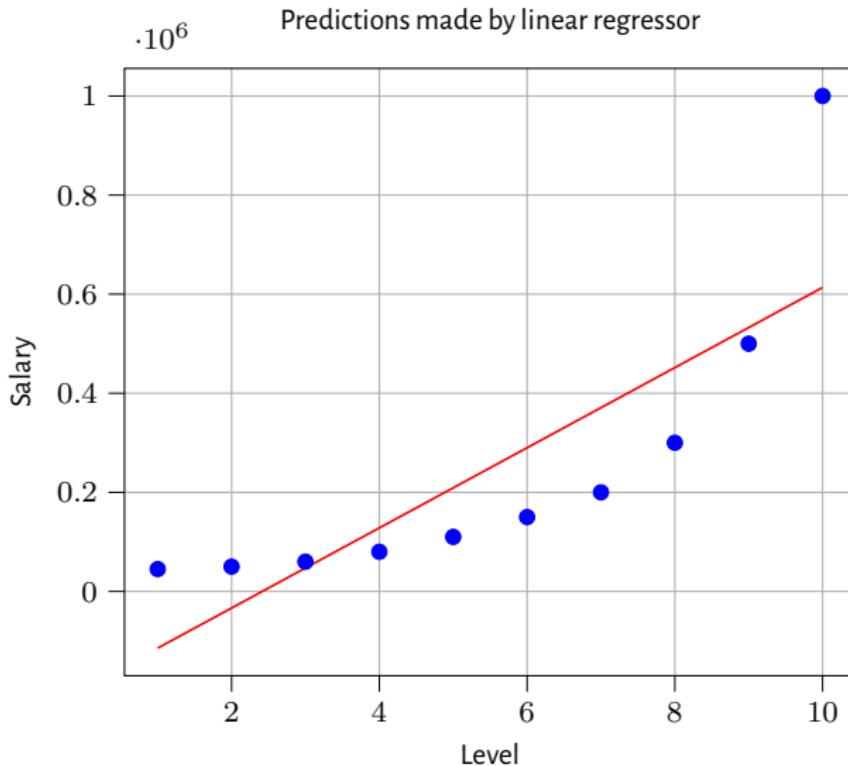
Visualising the linear regression predictions



Polynomial Regression (3/8)

```
[8]: plt.scatter(X, y, color = 'blue')
plt.plot(X, lr_1.predict(X), color = 'red')
plt.title('Predictions made by linear regressor')
plt.xlabel('Level')
plt.ylabel('Salary')
plt.grid()
```

Polynomial Regression (4/8)



Polynomial Regression (5/8)

```
[9]: from sklearn.preprocessing import PolynomialFeatures
```

```
[10]: poly_reg = PolynomialFeatures(degree=4)
X_poly = poly_reg.fit_transform(X)
print(X_poly[:5])
lr_2 = LinearRegression()
lr_2.fit(X_poly, y)
```

```
[[ 1.    1.    1.    1.    1.]
 [ 1.    2.    4.    8.   16.]
 [ 1.    3.    9.   27.   81.]
 [ 1.    4.   16.   64.  256.]
 [ 1.    5.   25.  125. 625.]]
```

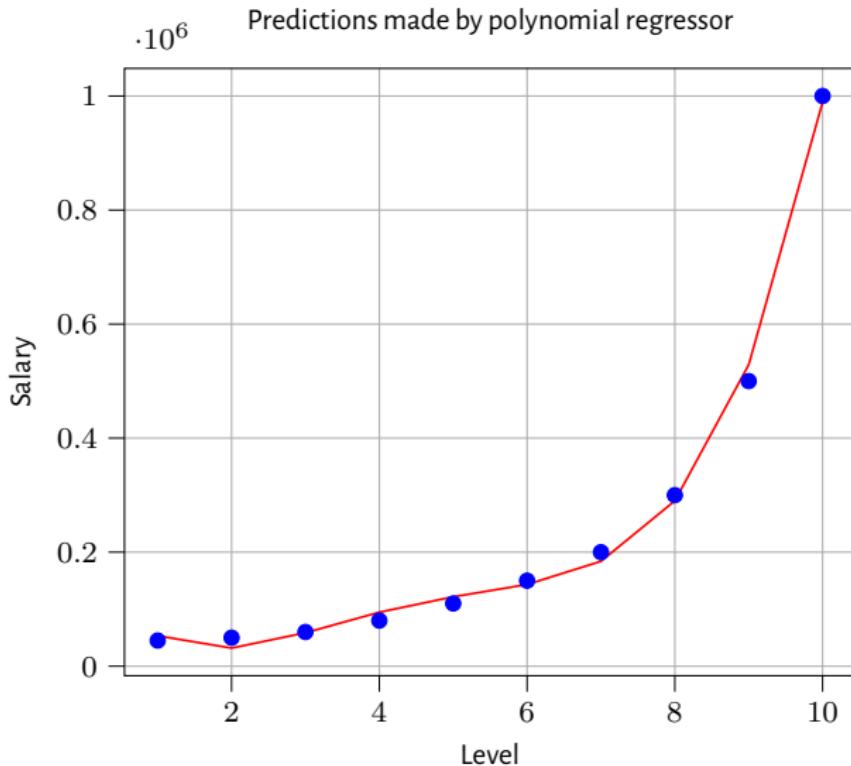
Polynomial Regression (6/8)

[10]: LinearRegression()

Visualising the polynomial regression predictions

```
[11]: plt.scatter(X, y, color='blue')
plt.plot(X, lr_2.predict(poly_reg.fit_transform(X)), color='red')
plt.title('Predictions made by polynomial regressor')
plt.xlabel('Level')
plt.ylabel('Salary')
plt.grid()
```

Polynomial Regression (7/8)



Polynomial Regression (8/8)

Predicting a new result using the linear regressor

```
[12]: lr_1.predict([[6.5]])
```

```
[12]: array([330378.78787879])
```

Predicting a new result using the polynomial regressor

```
[13]: lr_2.predict(poly_reg.fit_transform([[6.5]]))
```

```
[13]: array([158862.45265155])
```

Task #2

The yield y of a chemical process is a random variable whose value is considered to be a linear function of the temperature x . The following data of corresponding values of x and y is found:

Temperature in °C (x)	0	25	50	75	100
Yield in grams (y)	14	38	54	76	95

The linear regression model $y = \theta_0 + \theta_1 x$ is used. Determine the values of θ_0 , θ_1 .

- ① Using normal equation,
- ② Using gradient descent for 5 iterations.

$$y = \begin{bmatrix} 14 \\ 38 \\ 54 \\ 76 \\ 95 \end{bmatrix} \text{ et } X = \begin{bmatrix} 1 & 0 \\ 1 & 25 \\ 1 & 50 \\ 1 & 75 \\ 1 & 100 \end{bmatrix} \implies X^T X = \begin{bmatrix} 5 & 250 \\ 250 & 18750 \end{bmatrix}$$

$$\hat{\theta} = \begin{bmatrix} \hat{\theta}_0 \\ \hat{\theta}_1 \end{bmatrix} = \begin{bmatrix} 15.4 \\ 0.8 \end{bmatrix}$$

```
>>> import matplotlib.pyplot as plt
>>> import numpy as np

>>> X = np.array([[1,0], [1,25], [1,50], [1,75], [1,100]], dtype=np.float32)
>>> y = np.array([[14], [38], [54], [76], [95]])

>>> # NORMAL EQUATION
>>> XtX = X.T.dot(X)
>>> invXtX = np.linalg.inv(XtX)
>>> t_ne = invXtX.dot(np.matmul(X.T, y))

>>> print(t_ne)
[[15.39999944]
 [ 0.79999982]]

>>> X[:,1]= (X[:,1]-X[:,1].min())/X[:,1].max()
>>> y = (y-y.min())/y.max()

>>> # GRADIENT DESCENT
>>> t_gd = np.array([[1], [1]])
>>> alpha = .1
>>> vect = np.zeros(shape=(2, 1001))
>>> vect[:,0] = t_gd[[0,1],[0]]
>>> lost = []
```

```
>>> for k in range(1000):
...     eps = y-np.matmul(X, t_gd)
...     lost.append(1/(2*len(y))*eps.T.dot(eps)[[0],[0]][0])
...     t_gd = t_gd+alpha*1/len(y)*np.matmul((eps).T, X).T
...     vect[:,k+1] = t_gd[[0,1],[0]]
...
>>> print(vect[:, -1])
[0.01474565  0.84208938]

>>> plt.plot(vect[0, :], label=r'$\hat{\theta}_0$')
[<matplotlib.lines.Line2D object at 0x7fde144cdb80>]
>>> plt.plot(vect[1, :], label=r'$\hat{\theta}_1$')
[<matplotlib.lines.Line2D object at 0x7fde144cdf10>]
>>> plt.legend()
<matplotlib.legend.Legend object at 0x7fde144e41f0>
>>> plt.grid()
>>> plt.show()

>>> plt.plot(lost)
[<matplotlib.lines.Line2D object at 0x7fde000bcf40>]
>>> plt.grid()
>>> plt.show()
```

[Next...](#)

- 1 An overview
- 2 Supervised Learning
- 3 Unsupervised Learning
- 4 Deep Learning

Next...



- 1 An overview
- 2 Supervised Learning
- 3 Unsupervised Learning
- 4 Deep Learning

Further Reading



I. El Naqa and M. J. Murphy. "What Is Machine Learning?" In: *Machine Learning in Radiation Oncology: Theory and Applications*. Ed. by I. El Naqa, R. Li, and M. J. Murphy. Cham: Springer International Publishing, 2015, pp. 3–11. DOI: [10.1007/978-3-319-18305-3_1](https://doi.org/10.1007/978-3-319-18305-3_1).



J. Schmidt et al. "Recent advances and applications of machine learning in solid-state materials science". In: *npj Computational Materials* 5.1 (Aug. 2019). DOI: [10.1038/s41524-019-0221-0](https://doi.org/10.1038/s41524-019-0221-0) (cit. on p. 15).